



# AI reinvention made real

Accenture Innovation Challenge 2026

# Team details

TEAM NAME: 23je0841



## Sahil Bhavesh Choudhary

College: IIT (ISM) Dhanbad  
Stream: B.Tech - Petroleum Engineering  
Year of graduation: 2027



## Mayank Raj

College: IIT (ISM) Dhanbad  
Stream: B.Tech - Petroleum Engineering  
Year of graduation: 2027



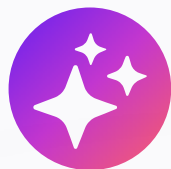
## Saket Kumar Singh

College: IIT (ISM) Dhanbad  
Stream: B.Tech - Petroleum Engineering  
Year of graduation: 2027

# ControlPlane.ai:

## Sidecar Orchestration

### The Problem: After the Fact Discovery



Current AI deployments lack real-time telemetry across inference planes. Waiting for user feedback to identify confidently wrong, expensive, or unsafe outputs is a critical operational failure.

### The Technical Approach



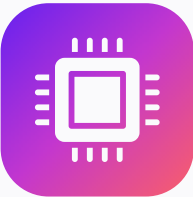
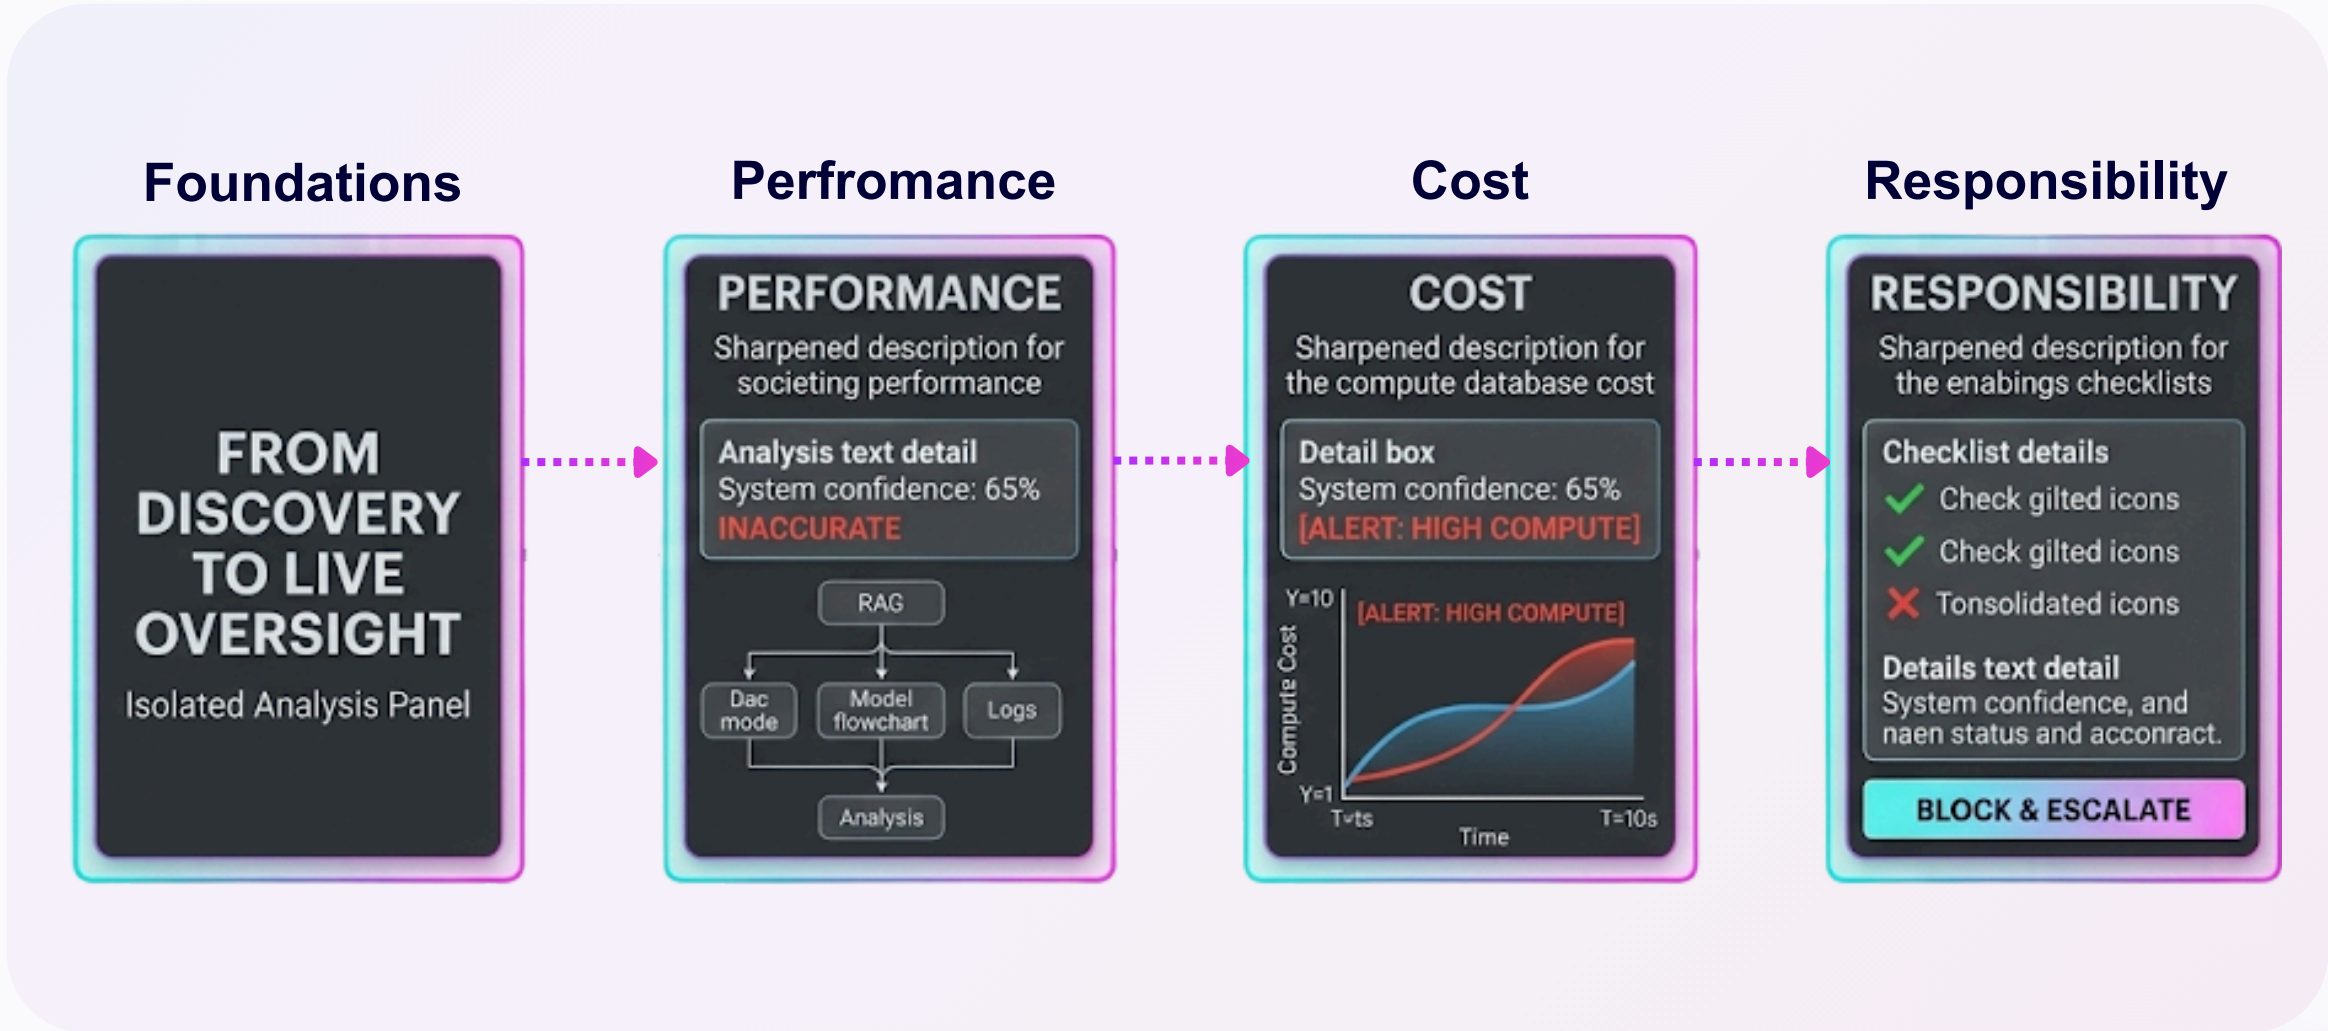
Using a non blocking sidecar proxy to intercept LLM payloads at the network layer. Leveraging an asynchronous, event-driven architecture (Kafka/Redis), it forks the stream; delivering the primary response to the user instantly while micro agents perform parallel, real time verification to eliminate system bottlenecks.



# Solution

## Asynchronous Evaluation Micro Agents

ControlPlane Checker decouples evaluation from generation. By utilizing an asynchronous, multi agent orchestration layer, we ensure continuous, real-time oversight across all dimensions without introducing system latency.



### Performance Agent:

Combats hallucinations using RAG verification and quantized Cross Encoders to score real-time semantic similarity against grounded enterprise context.



### Cost Telemetry:

Monitors active generation loops using streaming token heuristics to prevent compute burn, applying Redis caching and dynamic rate-limiting.

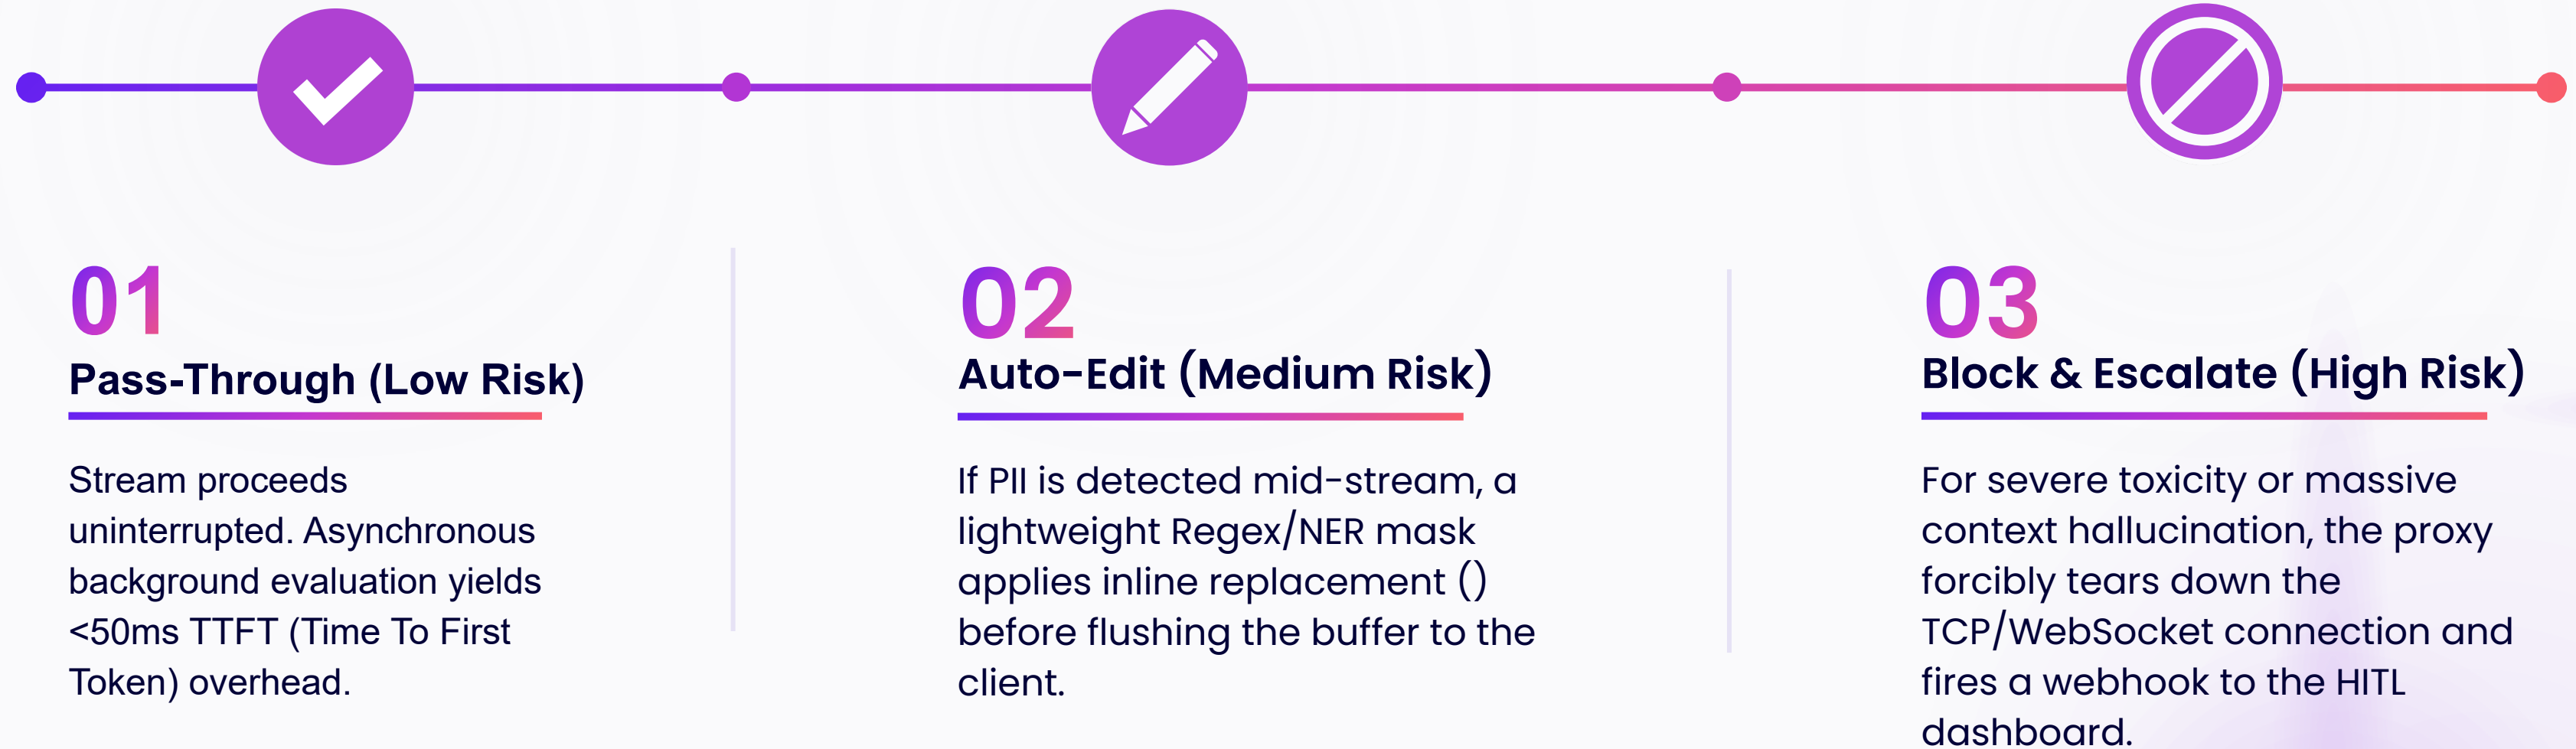


### Responsibility Layer:

Deploys ultra-fast NER pipelines and toxicity classifiers to scan outbound payloads, preventing data leakage and enforcing enterprise guardrails.

# Dynamic Action Engine & Mitigation

How the ControlPlane resolves latency and acts on flagged telemetry without crippling the primary LLM pipeline:



# Thank you

