



AI reinvention made real

Accenture Innovation Challenge 2026

Team details

TEAM NAME: 23je0841




Sahil Bhavesh Choudhary

College: IIT (ISM) Dhanbad
Stream: B.Tech - Petroleum Engineering
Year of graduation: 2027



Mayank Raj

College: IIT (ISM) Dhanbad
Stream: B.Tech - Petroleum Engineering
Year of graduation: 2027




Saket Kumar Singh

College: IIT (ISM) Dhanbad
Stream: B.Tech - Petroleum Engineering
Year of graduation: 2027

Repo Link: [Click Here](#)


Demo/Prototype Video link: [Click Here](#)




ControlPlane.ai:

Sidecar Orchestration

The Problem: After the Fact Discovery




Current AI deployments **lack real-time telemetry** across inference planes. Waiting for user feedback to identify confidently wrong, expensive, or unsafe outputs is a critical operational failure.




Technical Approach

Built on **Python** and **Starlette**, a non-blocking sidecar proxy instantly forwards LLM streams while parallel micro-agents evaluate risks asynchronously via a Redis event bus.

Architecture & Network Interception



Parallel Execution & Mitigation Flowchart



Competitive Matrix & Commercial Opportunity

Feature	Legacy Reactive Logging	In-Line Blocking Proxies	ControlPlane.ai
User Latency Impact	Zero (Post-generation)	High (Blocks generation)	Zero (<50ms TTFT)
Evaluation Timing	"After the Fact"	Synchronous / Bottlenecked	Asynchronous / Parallel
Mitigation Protocol	Reporting only	Binary (Block/Allow)	Pass, Auto-Edit, Block
Cost Telemetry	Delayed billing sync	Heuristic estimations	Real-time token burn

\$42.2B

Projected Global enterprise AI market by the end of 2026, while the Global footprint to **\$539.5B**.

15%

Estimated **15% (\$6.3B)** allocated to security, compliance, and observability.

\$3.5B

SAM for LLM - specific telemetry and guardrail globally.

45%

Growing at 45% CAGR as Gen AI moves from sandbox to Production.

A conservative early capture yields a **\$150M-\$200M** near term revenue opportunity.

Capturing the **\$3.5B SAM** for LLM telemetry by targeting the security and compliance budgets of **Fortune 500 AI deployments**.

Usage-based SaaS billed per million intercepted tokens, mirroring the cost structure of upstream providers like **OpenAI** and **Anthropic**.




Mitigation, Traction & Roadmap

Traction



- **Live Deployment:** Functional 5-service Docker Compose topology mimicking real enterprise deployment.
- **Core Proxy:** Streaming pass-through and mid-stream teardown fully operational.
- **Micro-Agents:** Performance, Cost, and Responsibility agents running 8 deterministic scenarios via an offline mock provider.

Roadmap



Phase 1	Week 1	Integrate fine-tuned neural cross-encoders.
Phase 2	Week 2	Add production vector stores & exact tokenizers.
Phase 3	Week 4	Launch SSO, auth, and multi-tenancy.

GTM



- **Developer-First Land & Expand:** Drive bottom-up adoption by targeting developers with the zero-key mock provider demo.
- **Enterprise Scaling:** Transition free developer usage into top-down, multi-tenant deployments mandated by internal security and DevOps teams.



Risk & Mitigation



01 Pass-Through (Low Risk)

Stream proceeds uninterrupted; background checks confirm safety. < 50ms TTFT overhead.



02 Auto-Edit (Medium Risk)

Luhn-validated Regex/NER masks PII mid-stream before flushing the buffer.



03 Block & Escalate (High Risk)

Proxy forcefully tears down the TCP/WebSocket connection mid-stream and fires a webhook to the live HITL Dashboard.

Thank you

Repo Link: [Click Here](#)

Demo/Prototype Video link: [Click Here](#)

Team 23je0841: Engineered by Saket, Mayank, and Sahil at IIT (ISM) Dhanbad.

